



```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
import plotly.express as px
import plotly.graph_objects as go
import plotly.io as pio
pio.templates.default = "plotly_white"
```

```
In [2]: df = pd.read_csv('tsla_2014_2023.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	date	open	high	low	close	volume	rsi_7	rsi_14	cci_7	cci_14	sma_50	ema_50	sma_100
0	2014-01-02	9.986667	10.165333	9.770000	10.006667	92826000	55.344071	54.440118	-37.373644	15.213422	9.682107	9.820167	10.494240
1	2014-01-03	10.000000	10.146000	9.906667	9.970667	70425000	53.742629	53.821521	-81.304471	17.481130	9.652800	9.826069	10.495693
2	2014-01-06	10.000000	10.026667	9.682667	9.800000	80416500	46.328174	50.870410	-123.427544	-37.824708	9.629467	9.825047	10.496740
3	2014-01-07	9.841333	10.026667	9.683333	9.957333	75511500	53.263037	53.406750	-84.784651	-20.779431	9.597747	9.830235	10.503407
4	2014-01-08	9.923333	10.246667	9.917333	10.085333	92448000	58.368660	55.423026	60.799662	43.570559	9.573240	9.840239	10.511147

```
In [4]: df.tail()
```

```
Out[4]:
```

	date	open	high	low	close	volume	rsi_7	rsi_14	cci_7	cci_14	sma_50	ema_50	sma_100
2511	2023-12-22	256.760010	258.220001	251.369995	252.539993	93249800	58.296612	58.137456	50.821325	80.672033	232.553000	240.3605	
2512	2023-12-26	254.490005	257.970001	252.910004	256.609985	86892400	63.570549	60.824035	93.909968	86.446838	232.662800	240.9978	
2513	2023-12-27	258.350006	263.339996	257.519989	261.440002	106494400	68.998630	63.793639	171.938770	119.554558	232.813200	241.7994	
2514	2023-12-28	263.660004	265.130005	252.710007	253.179993	113619900	53.186966	55.978816	45.772983	73.958135	232.779799	242.2457	
2515	2023-12-29	255.100006	255.190002	247.429993	248.479996	100615300	46.164227	52.070118	-98.880167	-0.320098	232.895800	242.4902	

```
In [5]: df.isnull().sum()
```

```
Out[5]:
```

date	0
open	0
high	0
low	0
close	0
volume	0
rsi_7	0
rsi_14	0
cci_7	0
cci_14	0
sma_50	0
ema_50	0
sma_100	0
ema_100	0
macd	0
bollinger	0
TrueRange	0
atr_7	0
atr_14	0
next_day_close	0
dtype:	int64

```
In [6]: df.describe()
```

Out[6]:

	open	high	low	close	volume	rsi_7	rsi_14	cci_7	cci_14	sma_50
count	2516.000000	2516.000000	2516.000000	2516.000000	2.516000e+03	2516.000000	2516.000000	2516.000000	2516.000000	2516.000000
mean	94.098510	96.172733	91.865096	94.072491	1.131986e+08	53.058382	52.862457	9.809933	13.202457	91.810735
std	108.593936	111.022486	105.911918	108.500301	7.547433e+07	18.239752	13.352063	100.975002	109.285239	106.581797
min	9.366667	9.800000	9.111333	9.289333	1.062000e+07	6.395305	16.564126	-233.333333	-297.930166	9.490973
25%	15.763167	16.082168	15.491167	15.814167	6.643185e+07	39.859440	43.595435	-76.876737	-78.543937	15.496080
50%	21.801001	22.198334	21.487666	21.877667	9.320775e+07	53.226417	51.621434	19.823624	24.702835	21.563733
75%	200.017505	204.525829	194.482498	200.049999	1.323710e+08	65.900330	61.937068	94.426550	99.180514	192.341650
max	411.470001	414.496674	405.666656	409.970001	9.140820e+08	97.460910	94.197983	233.333333	350.643337	357.870532

In [7]:

```
df.shape
```

Out[7]:

```
(2516, 20)
```

In [8]:

```
df.dtypes
```

Out[8]:

```
date                object
open               float64
high               float64
low                float64
close              float64
volume              int64
rsi_7               float64
rsi_14              float64
cci_7               float64
cci_14              float64
sma_50              float64
ema_50              float64
sma_100             float64
ema_100             float64
macd                float64
bollinger            float64
TrueRange            float64
atr_7                float64
atr_14               float64
next_day_close       float64
dtype: object
```

In [9]:

```
df['date'] = pd.to_datetime(df['date'])
df.set_index('date', inplace=True)
df.head()
```

Out[9]:

	open	high	low	close	volume	rsi_7	rsi_14	cci_7	cci_14	sma_50	ema_50	sma_100	er
date													
2014-01-02	9.986667	10.165333	9.770000	10.006667	92826000	55.344071	54.440118	-37.373644	15.213422	9.682107	9.820167	10.494240	9.
2014-01-03	10.000000	10.146000	9.906667	9.970667	70425000	53.742629	53.821521	-81.304471	17.481130	9.652800	9.826069	10.495693	9.
2014-01-06	10.000000	10.026667	9.682667	9.800000	80416500	46.328174	50.870410	-123.427544	-37.824708	9.629467	9.825047	10.496740	9.
2014-01-07	9.841333	10.026667	9.683333	9.957333	75511500	53.263037	53.406750	-84.784651	-20.779431	9.597747	9.830235	10.503407	9.
2014-01-08	9.923333	10.246667	9.917333	10.085333	92448000	58.368660	55.423026	60.799662	43.570559	9.573240	9.840239	10.511147	9.

In [10]:

```
import matplotlib.pyplot as plt
plt.figure(figsize=(20, 15))
plt.plot(df.index, df['open'], label='Open')
plt.plot(df.index, df['close'], label='Close')
plt.plot(df.index, df['high'], label='High')
plt.plot(df.index, df['low'], label='Low')
plt.legend()
```


Out[10]:

```
<matplotlib.legend.Legend at 0x26f19305190>
```



```
In [11]: df['high'].plot(figsize=(12,4))
```

```
Out[11]: <Axes: xlabel='date'>
```




```
In [12]: df['high'].head()
```

```
Out[12]: date
2014-01-02    10.165333
2014-01-03    10.146000
2014-01-06    10.026667
2014-01-07    10.026667
2014-01-08    10.246667
Name: high, dtype: float64
```


```
In [13]: df[['open','next_day_close']].plot(figsize=(12,5))
```

Out[13]: <Axes: xlabel='date'>



```
In [14]: df['high'].plot(xlim=['2017-01-01','2022-02-28'],ylim=[0,1000],figsize=(12,4),c='red')
```

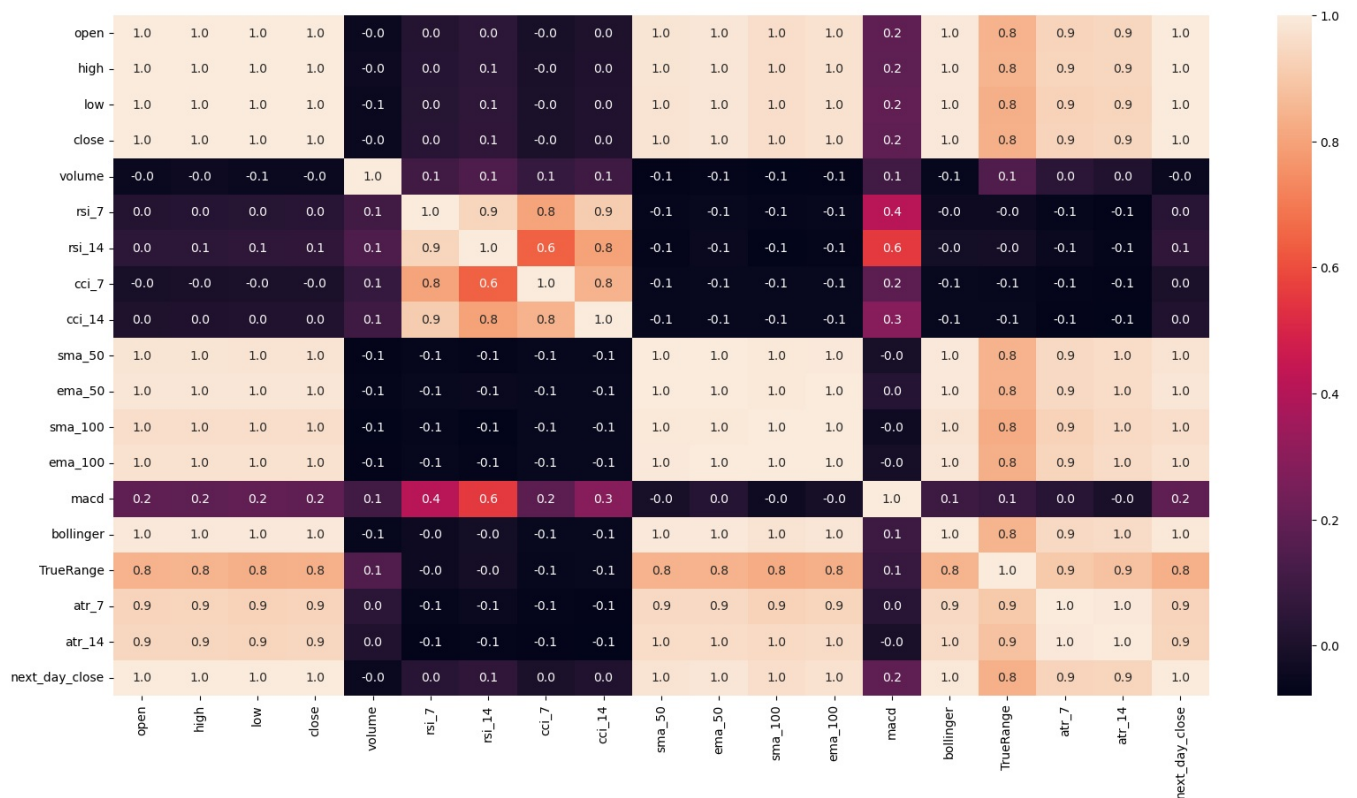
Out[14]: <Axes: xlabel='date'>



```
In [15]: df_corr = df.corr()

import seaborn as sns
plt.figure(figsize=(20, 10))
sns.heatmap(df_corr, annot=True, fmt='.1f')
```

Out[15]: <Axes: >



The “next_day_close” column has a 100% percent correlation with “open”, “close”, “low”, “high”, “sma_50”, and so on. This means if values in those columns increase the values in the “next_day_close” column will also increase.

```
In [16]: y = df['next_day_close']
X = df.drop(columns='next_day_close')
```

```
In [17]: X
```

Out[17]:

	open	high	low	close	volume	rsi_7	rsi_14	cci_7	cci_14	sma_50	ema_50	
date												
2014-01-02	9.986667	10.165333	9.770000	10.006667	92826000	55.344071	54.440118	-37.373644	15.213422	9.682107	9.820167	1
2014-01-03	10.000000	10.146000	9.906667	9.970667	70425000	53.742629	53.821521	-81.304471	17.481130	9.652800	9.826069	1
2014-01-06	10.000000	10.026667	9.682667	9.800000	80416500	46.328174	50.870410	-123.427544	-37.824708	9.629467	9.825047	1
2014-01-07	9.841333	10.026667	9.683333	9.957333	75511500	53.263037	53.406750	-84.784651	-20.779431	9.597747	9.830235	1
2014-01-08	9.923333	10.246667	9.917333	10.085333	92448000	58.368660	55.423026	60.799662	43.570559	9.573240	9.840239	1
...	
2023-12-22	256.760010	258.220001	251.369995	252.539993	93249800	58.296612	58.137456	50.821325	80.672033	232.553000	240.360582	24
2023-12-26	254.490005	257.970001	252.910004	256.609985	86892400	63.570549	60.824035	93.909968	86.446838	232.662800	240.997814	24
2023-12-27	258.350006	263.339996	257.519989	261.440002	106494400	68.998630	63.793639	171.938770	119.554558	232.813200	241.799468	24
2023-12-28	263.660004	265.130005	252.710007	253.179993	113619900	53.186966	55.978816	45.772983	73.958135	232.779799	242.245763	24
2023-12-29	255.100006	255.190002	247.429993	248.479996	100615300	46.164227	52.070118	-98.880167	-0.320098	232.895800	242.490243	24

2516 rows × 18 columns

In [18]:	y
Out[18]:	<pre> date 2014-01-02 9.970667 2014-01-03 9.800000 2014-01-06 9.957333 2014-01-07 10.085333 2014-01-08 9.835333 ... 2023-12-22 256.609985 2023-12-26 261.440002 2023-12-27 253.179993 2023-12-28 248.479996 2023-12-29 248.419998 Name: next_day_close, Length: 2516, dtype: float64 </pre>
In [19]:	<pre> from sklearn.model_selection import train_test_split X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=777) </pre>
In [20]:	<pre> from sklearn.preprocessing import StandardScaler scaler = StandardScaler() X_train_scaled = scaler.fit_transform(X_train) X_test_scaled = scaler.fit_transform(X_test) X_train_scaled.shape, X_test_scaled.shape # ((2012, 18), (504, 18)) </pre>
Out[20]:	((2012, 18), (504, 18))
In [21]:	<pre> from sklearn.linear_model import LinearRegression lr = LinearRegression() lr.fit(X_train, y_train) lr_prediction = lr.predict(X_test) from sklearn.metrics import mean_absolute_error lr_mae = mean_absolute_error(y_test, lr_prediction) </pre>
In [22]:	lr_mae
Out[22]:	2.3867493379877676

Our model is performing amazingly well. Our linear regression model is just giving a very low error.

2.38 means that, if the actual value is 12.38 our prediction would be \$10.

It's not that bad. Let's check other models' performance as well.

In [23]:	<pre> from sklearn.tree import DecisionTreeRegressor dtr = DecisionTreeRegressor() dtr.fit(X_train, y_train) dtr_prediction = dtr.predict(X_test) </pre>
----------	--

```
dtr_mae = mean_absolute_error(y_test, dtr_prediction)
dtr_mae
```

Out[23]: 3.588810936507936

```
In [24]: from sklearn.ensemble import RandomForestRegressor
rfr = RandomForestRegressor()
rfr.fit(X_train, y_train)
rfr_prediction = rfr.predict(X_test)
rfr_mae = mean_absolute_error(y_test, rfr_prediction)
rfr_mae
```

Out[24]: 2.662456560972218

```
In [25]: from sklearn.svm import SVR
svr = SVR()
svr.fit(X_train, y_train)
svr_prediction = svr.predict(X_test)
svr_mae = mean_absolute_error(y_test, svr_prediction)
svr_mae
```

Out[25]: 74.6380006751906

In []:

Processing math: 100%